

MSDS 422 ASSIGNMENT 1 – EXPLORING AND VISUALIZING DATA

Anaswar Jayakumar

Dr. Lawrence Fulton



Data preparation

The original data set contained negative values for cases which would affect how the case rate was calculated and in turn the log of the case rate since negative numbers aren't included in the domain of the log function. Likewise, there were also negative values in the deaths column of the original data set as well and such values needed to be taken care before proceeding to analyzing the data. I also needed to check the presence of null values in the original data set and remove them accordingly. This was especially important as null values could potentially interfere with later analysis. Specifically, I noticed the presence of null values in the popData2019 and the countryterritoryCode columns which would affect the calculation of the case and death rates. Min-max scaling was performed on the COVID-19 dataset in order to normalize the variables and standard scaling was performed in order to scale the variables to a standard range. Lastly, I created a time series for the cases, deaths, the cumulative number of cases, the case rate, the death rate, and the log of the case and death rates. When calculating the case rate and death rate, I used 100,000 as the denominator since the cumulative case count used 100,000 as well.

Data exploration/Data visualization

After preparing the dataset for analysis, I then analyzed the data to better understand the spread of COVID-19 as well as its fatality rate. To do this, I first created two additional columns, case_rate and death_rate and then took the log of the case_rate and death_rate columns. When calculating the log of the case and death rates, I added 0.1 since the log of 0 is undefined but case rate and death rate happened to have 0 as valid values. Adding a small constant ensures that the domain of the log is satisfied while not impacting the data significantly. Once, the case rate, death rate, log case rate and log death rate columns were added, I then generated the descriptive statistics of the updated COVID-19 dataset. To better visualize and understand the data, I generated a scatter plot modeling the case rate versus the death rate, histograms modeling the distribution of the case rate and death rate as well as the log of the case rate and death rate, time series graphs for the cases, deaths, the cumulative number of cases in the past 14 days, the case and death rate, and the log of the case and death rate, correlation matrices of the COVID-19 data and the time series, and boxplots of the case and death rates as well as the boxplots of the log of the case and death rates. The plots helped better understand the data as a whole as well as the spread of the disease and whether or not it is a predictor of the

fatality rate of the disease. Because the distribution of the case and death rate seemed exponential, it made sense to perform a log transformation on these columns. I chose log2.

Data scaling and comparisons

Min-max scaling was performed on the final COVID-19 dataset in order to normalize the variables and standard scaling was performed in order to scale the variables to a standard range. For both the min-max and standard scaling, I chose the columns cases, deaths, case_rate, death_rate, log_case_rate, and log_death_rate. Scaling methods such as the min-max and standard scaling methods preserve the overall shape of the distribution which is exponential in the case of cases, deaths, case_rate, and death_rate. In addition, min-max scaling was performed on the final COVID-19 dataset since the variables in the dataset are measured at different scales and therefore contribute unequally to the fitting of the model, thereby causing an inherent bias in the final model. One difference between the two methods is the range of the data and this is because the formula used in standard scaling involves dividing by the standard deviation as opposed to the min-max scaling which involves dividing by the range. The distributions generated by both methods are similar and comparable to the original distribution.

Insights from analysis

The descriptive statistics did provide some insight into the spread of COVID-19 as well as its fatality rate. I looked at the descriptive statistics of the original COVID-19 data, the descriptive statistics of the prepared COVID-19 data, and the descriptive statistics of the time series data. For the updated COVID-19 dataset, I noticed that the mean of the cases is larger than the upper quartile, the mean of the deaths is larger than the upper quartile as well, and that the standard deviation for the cases and deaths was significantly large as well. The boxplots of the case and death rates and the log of the case and death rates show most of the points outside of the 75th percentile. This makes sense as the distribution is exponential. The histogram confirms the shape of the distribution as well. This of course makes sense since cases of COVID-19 have experienced exponential growth and likewise, deaths from COVID-19 have also experienced exponential growth. For the prepared COVID-19 dataset, I generated a correlation matrix. I ignored columns such as day, month, and year and only looked at columns such as cases, deaths, total population, case rate and death rate. I expected case and deaths to be highly correlated and this was reflected in the correlation matrix, however the correlation between cases and population data was weaker than expected. Potential reasons include

countries experiencing an unequal number of cases and inaccuracy when reporting the actual number of cases and deaths. Even though the cases and deaths are highly correlated, the correlation between case rate and death rate seems to be weaker. Since the data was measured over a year, it made sense to model a time series for the cases, deaths, case rate, death rate, and the cumulative number of cases. The time series provided better insight into how the cases, deaths, case rate, death rate, and cumulative number of cases changes over time. The time series of the log of the death rate shows a peak followed by a decrease even though the log of the case rate stays relatively constant.

MSDS 422 - COVID 19 EDA (final)

April 4, 2021

1 Data Exploration, Data Preparation, and Data Visualization

Import all required Python libraries such as pandas, numpy, matplotlib, seaborn, etc

```
[1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.ensemble import ExtraTreesClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, roc_auc_score
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVC
```

Import COVID-19 data which is a CSV file

```
[2]: covid_19_dataset = '/Users/anaswarjayakumar/Desktop/Desktop - Anaswar's MacBook_
↳Pro/COVID-19 Data.csv'

covid_19_df = pd.read_csv(covid_19_dataset, sep = ",")

covid_19_df
```

```
[2]:
```

	dateRep	day	month	year	cases	deaths	countriesAndTerritories	\
0	14/12/2020	14	12	2020	746	6	Afghanistan	
1	13/12/2020	13	12	2020	298	9	Afghanistan	
2	12/12/20	12	12	2020	113	11	Afghanistan	
3	11/12/20	11	12	2020	63	10	Afghanistan	
4	10/12/20	10	12	2020	202	16	Afghanistan	
...
61895	31/12/2019	31	12	2019	0	0	United_States_of_America	
61896	24/03/2020	24	3	2020	0	1	Zimbabwe	
61897	31/12/2019	31	12	2019	0	0	Vietnam	

61898	22/03/2020	22	3	2020	1	0	Zimbabwe
61899	21/03/2020	21	3	2020	1	0	Zimbabwe

	geoId	countryterritoryCode	popData2019	continentExp	\
0	AF	AFG	38041757.0	Asia	
1	AF	AFG	38041757.0	Asia	
2	AF	AFG	38041757.0	Asia	
3	AF	AFG	38041757.0	Asia	
4	AF	AFG	38041757.0	Asia	
...	
61895	US	USA	329064917.0	America	
61896	ZW	ZWE	14645473.0	Africa	
61897	VN	VNM	96462108.0	Asia	
61898	ZW	ZWE	14645473.0	Africa	
61899	ZW	ZWE	14645473.0	Africa	

	Cumulative_number_for_14_days_of_COVID-19_cases_per_100000
0	9.013779
1	7.052776
2	6.868768
3	7.134266
4	6.968658
...	...
61895	NaN
61896	NaN
61897	NaN
61898	NaN
61899	NaN


[61900 rows x 12 columns]

Generate heat map for COVID-19 data frame

```
[3]: plt.figure(figsize=(15, 10))

sns.heatmap(covid_19_df.corr(), annot=True)
```

[3]: <AxesSubplot:>



Get descriptive statistics for COVID-19 data frame

```
[4]: covid_19_df.describe()
```

```
[4]:
```

	day	month	year	cases	deaths \
count	61900.000000	61900.000000	61900.000000	61900.000000	61900.000000
mean	15.628934	7.067157	2019.998918	1155.147237	26.055460
std	8.841582	2.954776	0.032882	6779.224479	131.227055
min	1.000000	1.000000	2019.000000	-8261.000000	-1918.000000

25%	8.000000	5.000000	2020.000000	0.000000	0.000000
50%	15.000000	7.000000	2020.000000	15.000000	0.000000
75%	23.000000	10.000000	2020.000000	273.000000	4.000000
max	31.000000	12.000000	2020.000000	234633.000000	4928.000000

```

popData2019 \
count 6.177700e+04
mean 4.098770e+07
std 1.531294e+08
min 8.150000e+02
25% 1.293120e+06
50% 7.169456e+06
75% 2.851583e+07
max 1.433784e+09

```


```

Cumulative_number_for_14_days_of_COVID-19_cases_per_100000
count 59021.000000
mean 66.320586
std 162.329240
min -147.419587
25% 0.757526
50% 6.724045
75% 52.572719
max 1900.836210

```


```
[5]: covid_19_df.hist('cases', bins = 100)
```

```
[5]: array([[<AxesSubplot:title={'center':'cases'}>]], dtype=object)
```

```
[6]: covid_19_df.hist('deaths', bins = 100)
```

```
[6]: array([[<AxesSubplot:title={'center':'deaths'}>]], dtype=object)
```



Convert the values in the geoId column to string and check which rows in the data frame contain null values in the popData2019 column

```
[7]: covid_19_df.geoId = covid_19_df.geoId.astype(str)

covid_19_df[covid_19_df['popData2019'].isnull()]
```

```
[7]:
```

	dateRep	day	month	year	cases	deaths	\
125	1/1/20	1	1	2020	0	0	
332	1/2/20	1	2	2020	0	0	
942	1/3/20	1	3	2020	0	0	
1838	1/11/20	1	11	2020	0	0	
1888	1/12/20	1	12	2020	0	0	
...	
60858	30/11/2020	30	11	2020	0	0	
60871	31/01/2020	31	1	2020	0	0	
60877	17/10/2020	17	10	2020	1	0	
61557	31/10/2020	31	10	2020	0	0	
61610	31/12/2019	31	12	2019	0	0	

	countriesAndTerritories	geoId	\
125	Cases_on_an_international_conveyance_Japan	JPG11668	
332	Cases_on_an_international_conveyance_Japan	JPG11668	
942	Cases_on_an_international_conveyance_Japan	JPG11668	
1838	Wallis_and_Futuna	WF	
1888	Wallis_and_Futuna	WF	
...	
60858	Wallis_and_Futuna	WF	
60871	Cases_on_an_international_conveyance_Japan	JPG11668	
60877	Wallis_and_Futuna	WF	
61557	Wallis_and_Futuna	WF	
61610	Cases_on_an_international_conveyance_Japan	JPG11668	

	countryterritoryCode	popData2019	continentExp	\
125	NaN	NaN	Other	
332	NaN	NaN	Other	
942	NaN	NaN	Other	
1838	NaN	NaN	Oceania	
1888	NaN	NaN	Oceania	
...	
60858	NaN	NaN	Oceania	
60871	NaN	NaN	Other	
60877	NaN	NaN	Oceania	
61557	NaN	NaN	Oceania	
61610	NaN	NaN	Other	

	Cumulative_number_for_14_days_of_COVID-19_cases_per_100000
125	NaN
332	NaN
942	NaN
1838	NaN
1888	NaN
...	...
60858	NaN
60871	NaN
60877	NaN
61557	NaN
61610	NaN

[123 rows x 12 columns]

Check which rows in the data frame contain null values in the countryterritoryCode column

```
[8]: covid_19_df[covid_19_df['countryterritoryCode'].isnull()]
```

```
[8]:
```

	dateRep	day	month	year	cases	deaths	\
125	1/1/20	1	1	2020	0	0	
332	1/2/20	1	2	2020	0	0	
942	1/3/20	1	3	2020	0	0	
1838	1/11/20	1	11	2020	0	0	
1888	1/12/20	1	12	2020	0	0	
...	
60858	30/11/2020	30	11	2020	0	0	
60871	31/01/2020	31	1	2020	0	0	
60877	17/10/2020	17	10	2020	1	0	
61557	31/10/2020	31	10	2020	0	0	
61610	31/12/2019	31	12	2019	0	0	

	countriesAndTerritories	geoId	\
125	Cases_on_an_international_conveyance_Japan	JPG11668	
332	Cases_on_an_international_conveyance_Japan	JPG11668	
942	Cases_on_an_international_conveyance_Japan	JPG11668	
1838	Wallis_and_Futuna	WF	
1888	Wallis_and_Futuna	WF	
...	
60858	Wallis_and_Futuna	WF	
60871	Cases_on_an_international_conveyance_Japan	JPG11668	
60877	Wallis_and_Futuna	WF	
61557	Wallis_and_Futuna	WF	
61610	Cases_on_an_international_conveyance_Japan	JPG11668	

	countryterritoryCode	popData2019	continentExp	\
125	NaN	NaN	Other	
332	NaN	NaN	Other	

942	NaN	NaN	Other
1838	NaN	NaN	Oceania
1888	NaN	NaN	Oceania
...
60858	NaN	NaN	Oceania
60871	NaN	NaN	Other
60877	NaN	NaN	Oceania
61557	NaN	NaN	Oceania
61610	NaN	NaN	Other

Cumulative_number_for_14_days_of_COVID-19_cases_per_100000	
125	NaN
332	NaN
942	NaN
1838	NaN
1888	NaN
...	...
60858	NaN
60871	NaN
60877	NaN
61557	NaN
61610	NaN

[123 rows x 12 columns]

Create a new COVID-19 data frame such that it doesn't contain any null values in the popData2019 and it doesn't contain any negative values in the cases and deaths columns. Having null and negative values will interfere with later analysis of the COVID-19 data frame

```
[9]: clean_covid_19_df = covid_19_df[~covid_19_df['popData2019'].isnull()]

clean_covid_19_df = clean_covid_19_df[(clean_covid_19_df['cases'] >= 0) &
    → (clean_covid_19_df['deaths'] >= 0)]

clean_covid_19_df
```

```
[9]:
```

	dateRep	day	month	year	cases	deaths	countriesAndTerritories \
0	14/12/2020	14	12	2020	746	6	Afghanistan
1	13/12/2020	13	12	2020	298	9	Afghanistan
2	12/12/20	12	12	2020	113	11	Afghanistan
3	11/12/20	11	12	2020	63	10	Afghanistan
4	10/12/20	10	12	2020	202	16	Afghanistan
...
61895	31/12/2019	31	12	2019	0	0	United_States_of_America
61896	24/03/2020	24	3	2020	0	1	Zimbabwe
61897	31/12/2019	31	12	2019	0	0	Vietnam
61898	22/03/2020	22	3	2020	1	0	Zimbabwe
61899	21/03/2020	21	3	2020	1	0	Zimbabwe

	geoId	countryterritoryCode	popData2019	continentExp	\
0	AF	AFG	38041757.0	Asia	
1	AF	AFG	38041757.0	Asia	
2	AF	AFG	38041757.0	Asia	
3	AF	AFG	38041757.0	Asia	
4	AF	AFG	38041757.0	Asia	
...	
61895	US	USA	329064917.0	America	
61896	ZW	ZWE	14645473.0	Africa	
61897	VN	VNM	96462108.0	Asia	
61898	ZW	ZWE	14645473.0	Africa	
61899	ZW	ZWE	14645473.0	Africa	

	Cumulative_number_for_14_days_of_COVID-19_cases_per_100000
0	9.013779
1	7.052776
2	6.868768
3	7.134266
4	6.968658
...	...
61895	NaN
61896	NaN
61897	NaN
61898	NaN
61899	NaN

[61753 rows x 12 columns]

Calculate the death rate and case rate and create two such columns in the new COVID-19 data frame. In addition calculate the log of the case rate and death rate and create two such columns as well. When calculating the log of the case and death rates, I added 0.1 in order to ensure that the log function doesnt deal with any potential negative numbers as negative numbers are not included in the domain of the log function

```
[10]: clean_covid_19_df['death_rate'] = clean_covid_19_df['deaths'] /
      ↪(clean_covid_19_df['popData2019'] / 100000)

clean_covid_19_df['case_rate'] = clean_covid_19_df['cases'] /
      ↪(clean_covid_19_df['popData2019'] / 100000)

clean_covid_19_df['log_case_rate'] = np.log2(clean_covid_19_df['case_rate'] + 0.
      ↪1)

clean_covid_19_df['log_death_rate'] = np.log2(clean_covid_19_df['death_rate'] +
      ↪0.1)
```

```
clean_covid_19_df
```

```
[10]:
```

	dateRep	day	month	year	cases	deaths	countriesAndTerritories	\
0	14/12/2020	14	12	2020	746	6	Afghanistan	
1	13/12/2020	13	12	2020	298	9	Afghanistan	
2	12/12/20	12	12	2020	113	11	Afghanistan	
3	11/12/20	11	12	2020	63	10	Afghanistan	
4	10/12/20	10	12	2020	202	16	Afghanistan	
...	
61895	31/12/2019	31	12	2019	0	0	United_States_of_America	
61896	24/03/2020	24	3	2020	0	1	Zimbabwe	
61897	31/12/2019	31	12	2019	0	0	Vietnam	
61898	22/03/2020	22	3	2020	1	0	Zimbabwe	
61899	21/03/2020	21	3	2020	1	0	Zimbabwe	

	geoId	countryterritoryCode	popData2019	continentExp	\
0	AF	AFG	38041757.0	Asia	
1	AF	AFG	38041757.0	Asia	
2	AF	AFG	38041757.0	Asia	
3	AF	AFG	38041757.0	Asia	
4	AF	AFG	38041757.0	Asia	
...	
61895	US	USA	329064917.0	America	
61896	ZW	ZWE	14645473.0	Africa	
61897	VN	VNM	96462108.0	Asia	
61898	ZW	ZWE	14645473.0	Africa	
61899	ZW	ZWE	14645473.0	Africa	

	Cumulative_number_for_14_days_of_COVID-19_cases_per_100000	death_rate	\
0	9.013779	0.015772	
1	7.052776	0.023658	
2	6.868768	0.028916	
3	7.134266	0.026287	
4	6.968658	0.042059	
...	
61895	NaN	0.000000	
61896	NaN	0.006828	
61897	NaN	0.000000	
61898	NaN	0.000000	
61899	NaN	0.000000	

	case_rate	log_case_rate	log_death_rate
0	1.961003	1.043347	-3.110640
1	0.783350	-0.178943	-3.015570
2	0.297042	-1.332636	-2.955501
3	0.165607	-1.912632	-2.985223
4	0.530995	-0.664298	-2.815437

```

...
61895  0.000000  -3.321928  -3.321928
61896  0.000000  -3.321928  -3.226638
61897  0.000000  -3.321928  -3.321928
61898  0.006828  -3.226638  -3.321928
61899  0.006828  -3.226638  -3.321928

```

[61753 rows x 16 columns]

Get the descriptive statistics of the new COVID-19 data frame

```
[11]: clean_covid_19_df.describe()
```

```

[11]:
      count  day      month      year      cases      deaths \
count  61753.000000  61753.000000  61753.000000  61753.000000  61753.000000
mean    15.629945    7.069227    2019.998931    1158.071689    26.083607
std      8.841257    2.950149     0.032675    6786.916211    130.238403
min       1.000000    1.000000    2019.000000     0.000000     0.000000
25%       8.000000    5.000000    2020.000000     0.000000     0.000000
50%      15.000000    7.000000    2020.000000    16.000000     0.000000
75%      23.000000   10.000000    2020.000000    276.000000     4.000000
max      31.000000   12.000000    2020.000000  234633.000000   4928.000000

```

```

      popData2019 \
count  6.175300e+04
mean    4.099461e+07
std     1.531581e+08
min     8.150000e+02
25%     1.293120e+06
50%     7.169456e+06
75%     2.851583e+07
max     1.433784e+09

```

```

      Cumulative_number_for_14_days_of_COVID-19_cases_per_100000 \
count  58997.000000
mean    66.329444
std     162.354715
min    -147.419587
25%      0.757526
50%      6.724045
75%     52.559960
max     1900.836210

```

```


      count  death_rate  case_rate  log_case_rate  log_death_rate
count  61753.000000  61753.000000  61753.000000  61753.000000
mean      0.081945    4.847870    -0.601635    -2.889706
std      0.312673   14.933306     2.841068     0.849840
min      0.000000    0.000000    -3.321928    -3.321928

```

25%	0.000000	0.000000	-3.321928	-3.321928
50%	0.000000	0.260533	-1.471798	-3.321928
75%	0.034235	3.124746	1.689186	-2.897167
max	20.036065	858.895706	9.746507	4.331710

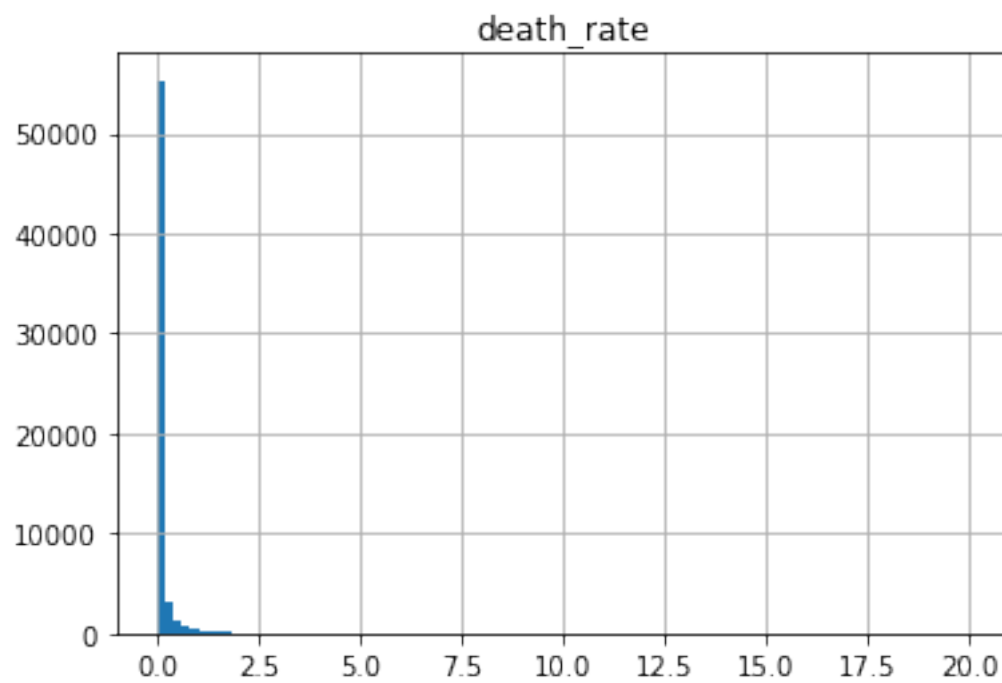
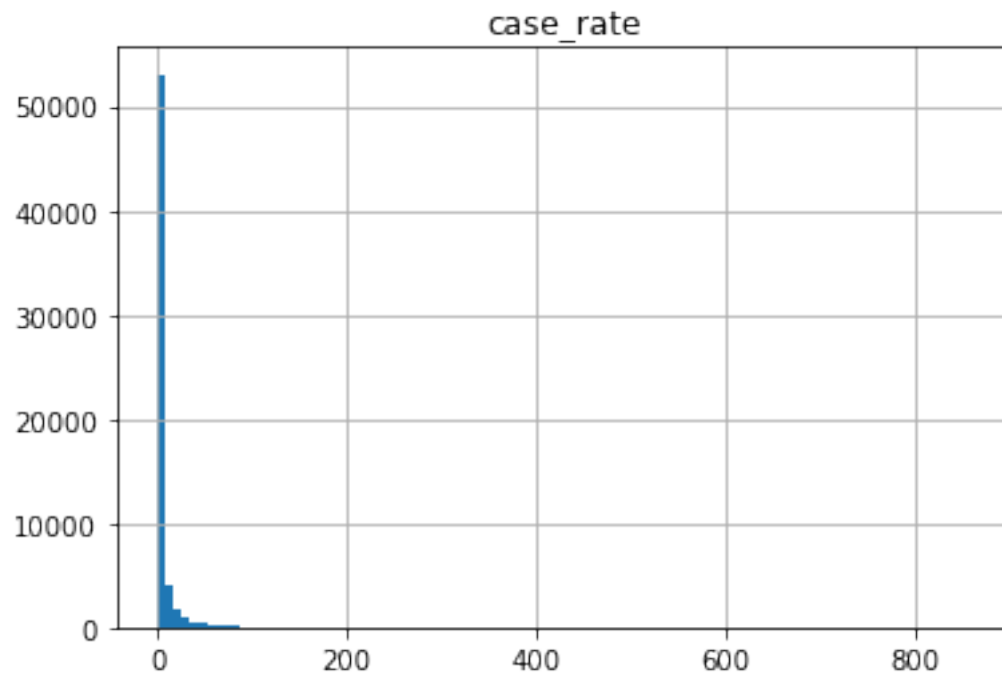
Generate a scatter plot of the updated COVID-19 data frame where the case rate is on the x-axis and the death rate is on the y-axis

```
[12]: covid_19_df_plot = clean_covid_19_df.plot.scatter(x = 'case_rate', y = 'death_rate', c='DarkBlue')
```



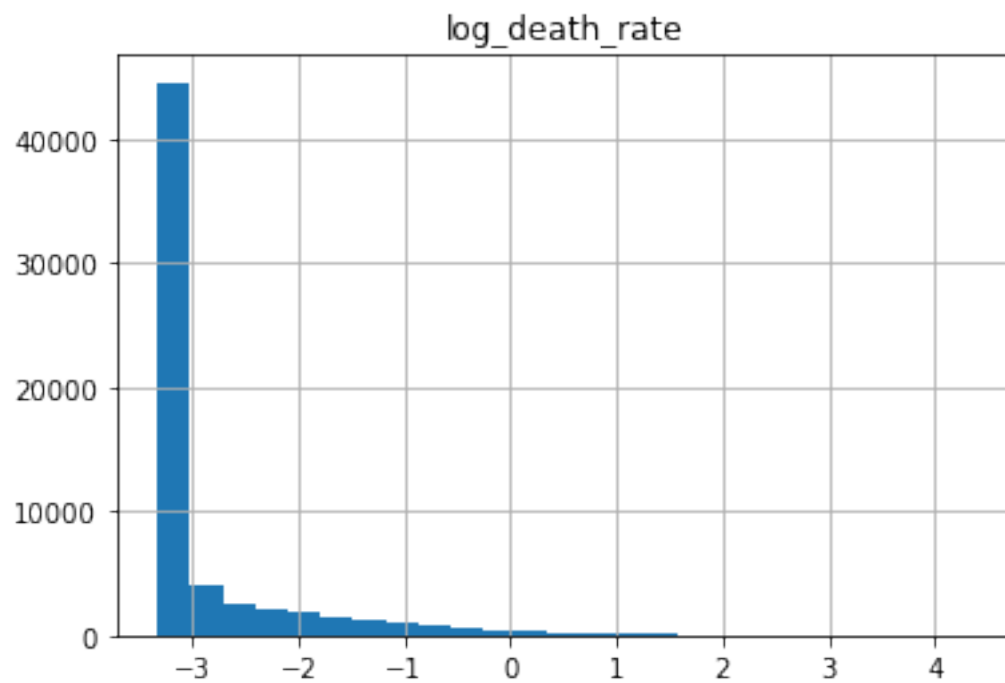

Generate histograms for the case rate and death rate.

```
[13]: covid_19_df_hist = clean_covid_19_df.hist('case_rate', bins = 100)
covid_19_df_hist = clean_covid_19_df.hist('death_rate', bins = 100)
```

Generate histograms for the log of the case and death rates.

```
[14]: covid_19_df_hist = clean_covid_19_df.hist('log_case_rate', bins = 25)
covid_19_df_hist = clean_covid_19_df.hist('log_death_rate', bins = 25)
```



Generate a correlation matrix for the updated COVID-19 data frame.

```
[15]: plt.figure(figsize=(15, 10))


sns.heatmap(clean_covid_19_df.corr(), annot=True)
```

[15]: <AxesSubplot:>

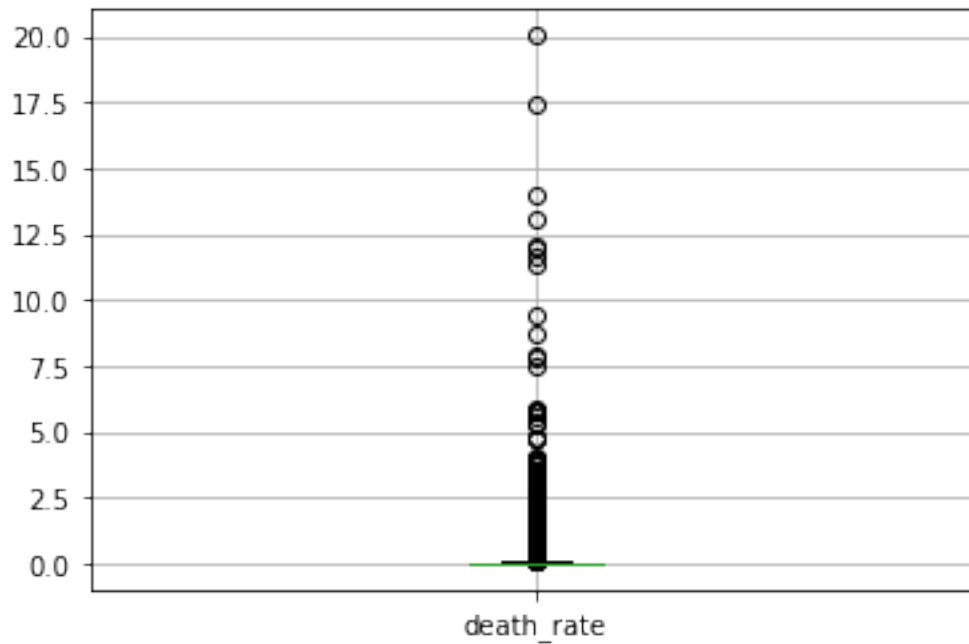


Generate the boxplots for the case rate and the death rate

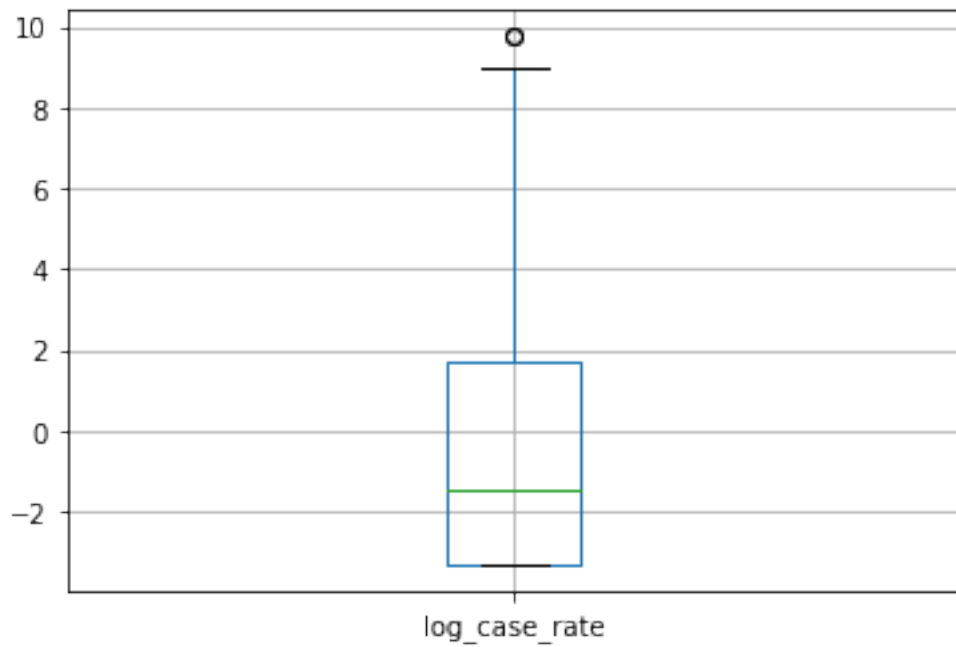
```
[16]: covid_19_df_box_plot_case_rate = clean_covid_19_df.boxplot('case_rate')
```



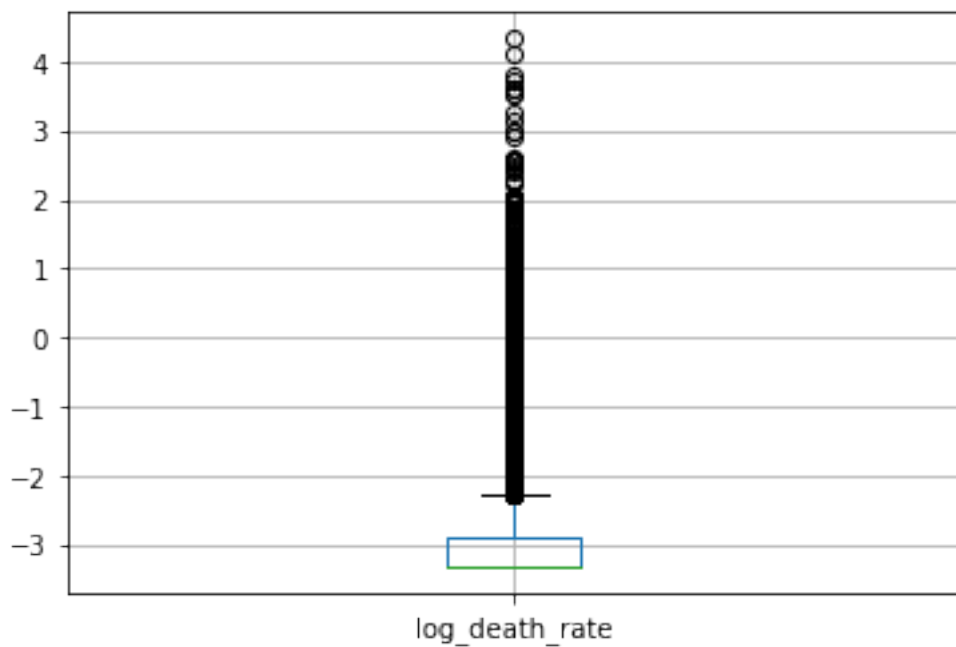
```
[17]: covid_19_df_box_plot_death_rate = clean_covid_19_df.boxplot('death_rate')
```



```
[18]: covid_19_df_box_plot_log_case_rate = clean_covid_19_df.boxplot('log_case_rate')
```



```
[19]: covid_19_df_box_plot_log_death_rate = clean_covid_19_df.  
      ↪boxplot('log_death_rate')
```



2 Min-max normalization and standard scaling for COVID-19 data

```
[20]: # Import MinMaxScaler from sklearn
from sklearn.preprocessing import MinMaxScaler

# build scalar model
scaler = MinMaxScaler()

# Perform min-max scaling
clean_covid_19_df_min_max = scaler.fit_transform(clean_covid_19_df[['cases',
    ↪ 'deaths', 'death_rate', 'case_rate',
    ↪
    ↪ 'log_death_rate', 'log_case_rate']])

# Verify min and max value
print(clean_covid_19_df_min_max.min(axis = 0))
print(clean_covid_19_df_min_max.max(axis = 0))
```

```
[0. 0. 0. 0. 0. 0.]
```

```
[1. 1. 1. 1. 1. 1.]
```

```
[21]: # Effect of Min-Max normalization in a visual example

fig, axes = plt.subplots(1,2)

axes[0].scatter(clean_covid_19_df['cases'], clean_covid_19_df['deaths'],
    ↪ c='darkblue')
axes[0].set_title("Original Cases vs Deaths")

axes[1].scatter(clean_covid_19_df_min_max[:,0], clean_covid_19_df_min_max[:,1],
    ↪ c='y')
axes[1].set_title("MinMax scaled Cases vs Deaths")

plt.show()

fig, axes = plt.subplots(1,2)

axes[0].scatter(clean_covid_19_df['case_rate'],
    ↪ clean_covid_19_df['death_rate'], c='darkblue')
axes[0].set_title("Original CR vs DR")

axes[1].scatter(clean_covid_19_df_min_max[:,3], clean_covid_19_df_min_max[:,2],
    ↪ c='y')
axes[1].set_title("MinMax scaled CR vs DR")

plt.show()
```

```


fig, axes = plt.subplots(1,2)

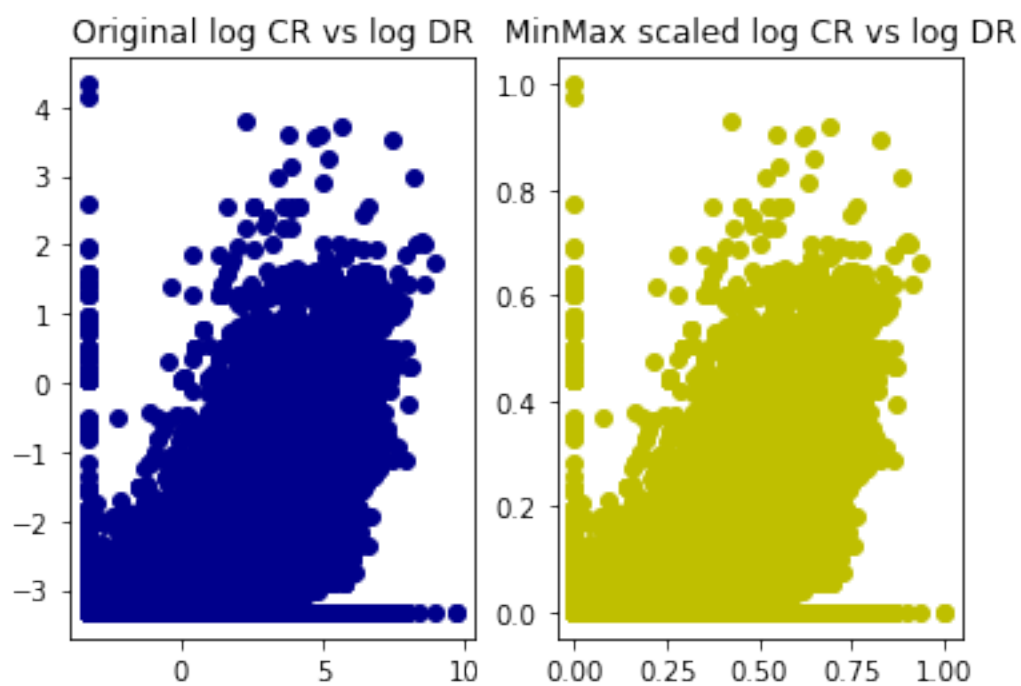
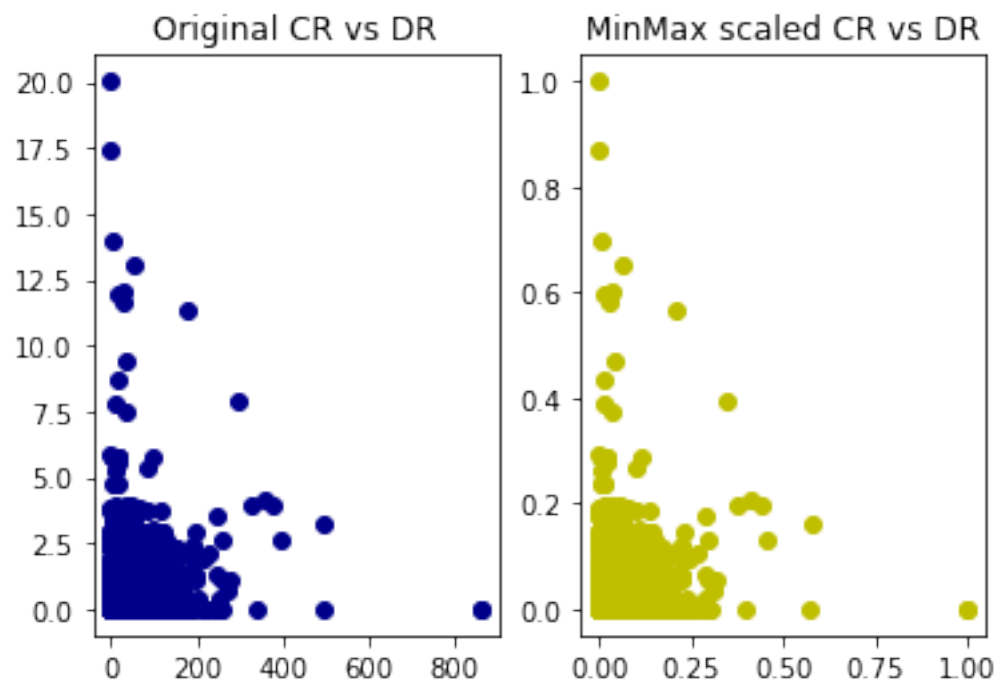
axes[0].scatter(clean_covid_19_df['log_case_rate'],
                ↪clean_covid_19_df['log_death_rate'], c='darkblue')
axes[0].set_title("Original log CR vs log DR")

axes[1].scatter(clean_covid_19_df_min_max[:,5], clean_covid_19_df_min_max[:,4],
                ↪c='y')
axes[1].set_title("MinMax scaled log CR vs log DR")

plt.show()

```






```
[22]: # Import StandardScaler from sklearn
from sklearn.preprocessing import StandardScaler

# build scalar model
scaler = StandardScaler()

std_scaled = scaler.fit_transform(clean_covid_19_df[['cases', 'deaths',
↳ 'death_rate', 'case_rate',
                                                    'log_death_rate',
↳ 'log_case_rate']])

print(std_scaled)

print("\n")

# Verify min and max value
print(std_scaled.min(axis = 0))
print(std_scaled.max(axis = 0))
```

```
[[-0.06071609 -0.15420774 -0.21163944 -0.19331889 -0.25997379  0.57900605]
 [-0.12672599 -0.13117287 -0.18641775 -0.27218038 -0.14810481  0.14878049]
 [-0.15398454 -0.1158163  -0.16960329 -0.30474596 -0.07742182 -0.25730011]
 ...
 [-0.17063435 -0.20027747 -0.26208282 -0.32463736 -0.50859702 -0.95749739]
 [-0.17048701 -0.20027747 -0.26208282 -0.32418012 -0.50859702 -0.92395674]
 [-0.17048701 -0.20027747 -0.26208282 -0.32418012 -0.50859702 -0.92395674]]

[-0.17063435 -0.20027747 -0.26208282 -0.32463736 -0.50859702 -0.95749739]
[34.40101682 37.63833061 63.8184198  57.19127212  8.49745551  3.6423725 ]
```

```
[23]: # Effect of Standard Scaling in a visual example

fig, axes = plt.subplots(1,2)

axes[0].scatter(clean_covid_19_df['cases'], clean_covid_19_df['deaths'],
↳ c='darkblue')
axes[0].set_title("Original Cases vs Deaths")

axes[1].scatter(std_scaled[:,0], std_scaled[:,1], c='y')
axes[1].set_title("Standard scaled Cases vs Deaths")

plt.show()

fig, axes = plt.subplots(1,2)
```

```

axes[0].scatter(clean_covid_19_df['case_rate'],
                ↪clean_covid_19_df['death_rate'], c='darkblue')
axes[0].set_title("Original CR vs DR")

axes[1].scatter(std_scaled[:,3], std_scaled[:,2], c='y')
axes[1].set_title("Standard scaled CR vs DR")

plt.show()


fig, axes = plt.subplots(1,2)



axes[0].scatter(clean_covid_19_df['log_case_rate'],
                ↪clean_covid_19_df['log_death_rate'], c='darkblue')
axes[0].set_title("Original log CR vs log DR")

axes[1].scatter(std_scaled[:,5], std_scaled[:,4], c='y')
axes[1].set_title("Standard scaled log CR vs log DR")

plt.show()

```





3 Time Series for updated Covid-19 data frame

[24]: *# Create datetime index for time series*

```
clean_covid_19_df.index = pd.to_datetime(clean_covid_19_df['dateRep'], dayfirst=  
    ↪= True)  
clean_covid_19_df
```

[24]:

	dateRep	day	month	year	cases	deaths	\
--	---------	-----	-------	------	-------	--------	---

dateRep							
2020-12-14	14/12/2020	14	12	2020	746	6	
2020-12-13	13/12/2020	13	12	2020	298	9	
2020-12-12	12/12/20	12	12	2020	113	11	
2020-12-11	11/12/20	11	12	2020	63	10	
2020-12-10	10/12/20	10	12	2020	202	16	
...			
2019-12-31	31/12/2019	31	12	2019	0	0	
2020-03-24	24/03/2020	24	3	2020	0	1	
2019-12-31	31/12/2019	31	12	2019	0	0	
2020-03-22	22/03/2020	22	3	2020	1	0	
2020-03-21	21/03/2020	21	3	2020	1	0	

	countriesAndTerritories	geoId	countryterritoryCode	popData2019	\
--	-------------------------	-------	----------------------	-------------	---

dateRep					
2020-12-14	Afghanistan	AF	AFG	38041757.0	
2020-12-13	Afghanistan	AF	AFG	38041757.0	
2020-12-12	Afghanistan	AF	AFG	38041757.0	
2020-12-11	Afghanistan	AF	AFG	38041757.0	
2020-12-10	Afghanistan	AF	AFG	38041757.0	
...	
2019-12-31	United_States_of_America	US	USA	329064917.0	
2020-03-24	Zimbabwe	ZW	ZWE	14645473.0	
2019-12-31	Vietnam	VN	VNM	96462108.0	
2020-03-22	Zimbabwe	ZW	ZWE	14645473.0	
2020-03-21	Zimbabwe	ZW	ZWE	14645473.0	

	continentExp	\
--	--------------	---

dateRep		
2020-12-14	Asia	
2020-12-13	Asia	
2020-12-12	Asia	
2020-12-11	Asia	
2020-12-10	Asia	
...	...	
2019-12-31	America	
2020-03-24	Africa	
2019-12-31	Asia	

```
2020-03-22      Africa
2020-03-21      Africa
```

```

Cumulative_number_for_14_days_of_COVID-19_cases_per_100000  \
dateRep
2020-12-14      9.013779
2020-12-13      7.052776
2020-12-12      6.868768
2020-12-11      7.134266
2020-12-10      6.968658
...
2019-12-31      NaN
2020-03-24      NaN
2019-12-31      NaN
2020-03-22      NaN
2020-03-21      NaN
```

```

death_rate  case_rate  log_case_rate  log_death_rate
dateRep
2020-12-14    0.015772    1.961003      1.043347      -3.110640
2020-12-13    0.023658    0.783350     -0.178943      -3.015570
2020-12-12    0.028916    0.297042     -1.332636      -2.955501
2020-12-11    0.026287    0.165607     -1.912632      -2.985223
2020-12-10    0.042059    0.530995     -0.664298      -2.815437
...
2019-12-31      0.000000    0.000000     -3.321928      -3.321928
2020-03-24    0.006828    0.000000     -3.321928      -3.226638
2019-12-31    0.000000    0.000000     -3.321928      -3.321928
2020-03-22    0.000000    0.006828     -3.226638      -3.321928
2020-03-21    0.000000    0.006828     -3.226638      -3.321928
```

```
[61753 rows x 16 columns]
```

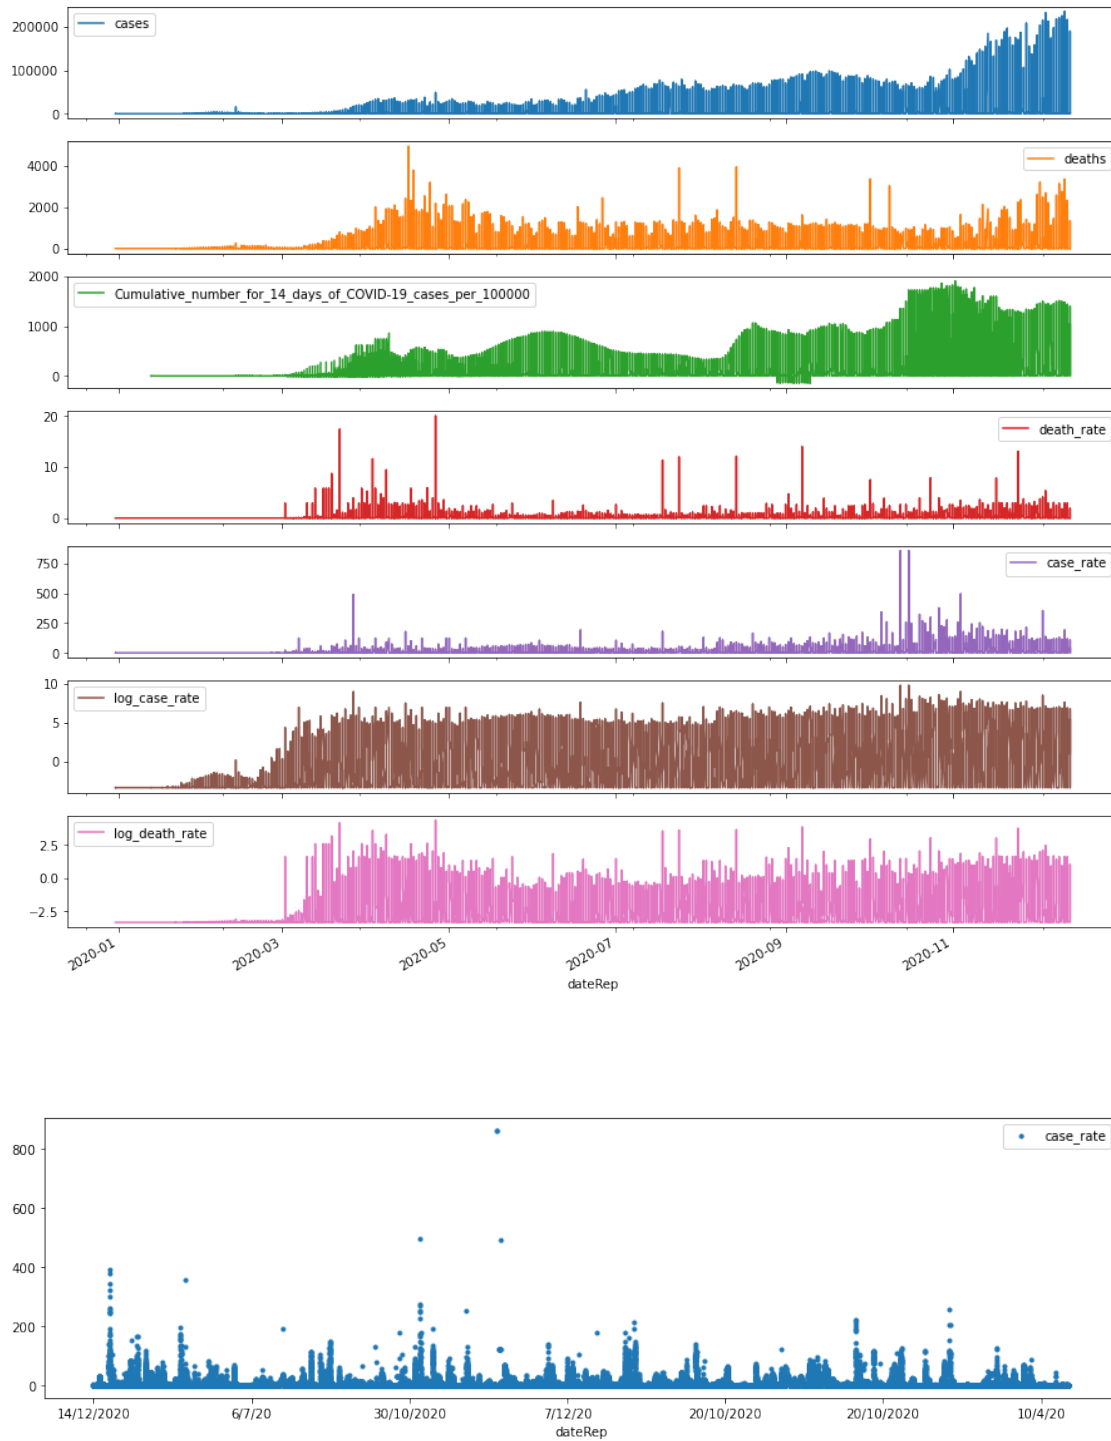
Generate time series for the cases, deaths, Cumulative_number_for_14_days_of_COVID-19_cases_per_100000, death_rate, case_rate, log_case_rate, and log_death_rate columns. In addition, generate a time series modeling the case rate over time

```
[25]: time_series_covid_19_df = clean_covid_19_df.drop(columns = ["popData2019",
    ↪ "day", "month", "year"])

time_series_covid_19_df.plot(subplots=True, figsize=(15,15))

time_series_covid_19_df.plot(x = 'dateRep', y = 'case_rate', style='.',
    ↪ figsize=(15,4))
```

```
[25]: <AxesSubplot:xlabel='dateRep'>
```



```
[26]: time_series_covid_19_df
```

[26]:

	dateRep	cases	deaths	countriesAndTerritories	geoId	\
dateRep						
2020-12-14	14/12/2020	746	6	Afghanistan	AF	
2020-12-13	13/12/2020	298	9	Afghanistan	AF	
2020-12-12	12/12/20	113	11	Afghanistan	AF	
2020-12-11	11/12/20	63	10	Afghanistan	AF	
2020-12-10	10/12/20	202	16	Afghanistan	AF	
...	
2019-12-31	31/12/2019	0	0	United_States_of_America	US	
2020-03-24	24/03/2020	0	1	Zimbabwe	ZW	
2019-12-31	31/12/2019	0	0	Vietnam	VN	
2020-03-22	22/03/2020	1	0	Zimbabwe	ZW	
2020-03-21	21/03/2020	1	0	Zimbabwe	ZW	

	countryterritoryCode	continentExp	\
dateRep			
2020-12-14	AFG	Asia	
2020-12-13	AFG	Asia	
2020-12-12	AFG	Asia	
2020-12-11	AFG	Asia	
2020-12-10	AFG	Asia	
...	
2019-12-31	USA	America	
2020-03-24	ZWE	Africa	
2019-12-31	VNM	Asia	
2020-03-22	ZWE	Africa	
2020-03-21	ZWE	Africa	

	Cumulative_number_for_14_days_of_COVID-19_cases_per_100000	\
dateRep		
2020-12-14	9.013779	
2020-12-13	7.052776	
2020-12-12	6.868768	
2020-12-11	7.134266	
2020-12-10	6.968658	
...	...	
2019-12-31	NaN	
2020-03-24	NaN	
2019-12-31	NaN	
2020-03-22	NaN	
2020-03-21	NaN	

	death_rate	case_rate	log_case_rate	log_death_rate
dateRep				
2020-12-14	0.015772	1.961003	1.043347	-3.110640
2020-12-13	0.023658	0.783350	-0.178943	-3.015570
2020-12-12	0.028916	0.297042	-1.332636	-2.955501

2020-12-11	0.026287	0.165607	-1.912632	-2.985223
2020-12-10	0.042059	0.530995	-0.664298	-2.815437
...
2019-12-31	0.000000	0.000000	-3.321928	-3.321928
2020-03-24	0.006828	0.000000	-3.321928	-3.226638
2019-12-31	0.000000	0.000000	-3.321928	-3.321928
2020-03-22	0.000000	0.006828	-3.226638	-3.321928
2020-03-21	0.000000	0.006828	-3.226638	-3.321928


[61753 rows x 12 columns]

Correlation Matrix for Time Series of COVID-19 data

```
[27]: plt.figure(figsize=(15, 10))

sns.heatmap(time_series_covid_19_df.corr(), annot=True)
```

[27]: <AxesSubplot:>




```
[28]: # Get descriptive statistics of the time series
time_series_covid_19_df.describe()
```

```
[28]:
```

	cases	deaths \
count	61753.000000	61753.000000
mean	1158.071689	26.083607
std	6786.916211	130.238403
min	0.000000	0.000000
25%	0.000000	0.000000
50%	16.000000	0.000000
75%	276.000000	4.000000
max	234633.000000	4928.000000

	Cumulative_number_for_14_days_of_COVID-19_cases_per_100000 \
count	58997.000000
mean	66.329444
std	162.354715
min	-147.419587
25%	0.757526
50%	6.724045
75%	52.559960
max	1900.836210

	death_rate	case_rate	log_case_rate	log_death_rate
count	61753.000000	61753.000000	61753.000000	61753.000000
mean	0.081945	4.847870	-0.601635	-2.889706
std	0.312673	14.933306	2.841068	0.849840
min	0.000000	0.000000	-3.321928	-3.321928
25%	0.000000	0.000000	-3.321928	-3.321928
50%	0.000000	0.260533	-1.471798	-3.321928
75%	0.034235	3.124746	1.689186	-2.897167
max	20.036065	858.895706	9.746507	4.331710

```
[ ]:
```